

# Statistical Learning — Practical Exam

Gabriele Socrate

2025-11-20

## Exercises

### Exercise 1

```
load("data_202511201120.RData")
names(data.all)

## [1] "Y"                  "age"                "sex"                "baseline_sbp"
## [5] "BMI"                "smoker"              "exercise_level"   "salt_intake"
## [9] "treatment"          "adherence"           "diabetes"

num.cov = data.all[, -c(1, 3, 6, 7, 9, 11)]
num.cov.scaled = scale(num.cov)
summary(num.cov.scaled)

##      age        baseline_sbp        BMI        salt_intake
##  Min. :-1.74088  Min. :-3.1830  Min. :-2.81839  Min. :-3.453407
##  1st Qu.:-0.84030  1st Qu.:-0.6518  1st Qu.:-0.63045  1st Qu.:-0.657471
##  Median :-0.03251  Median :-0.0177  Median :-0.03243  Median : 0.000434
##  Mean   : 0.00000  Mean   : 0.0000  Mean   : 0.00000  Mean   : 0.000000
##  3rd Qu.: 0.88970  3rd Qu.: 0.6561  3rd Qu.: 0.71888  3rd Qu.: 0.695864
##  Max.   : 1.77533  Max.   : 3.2019  Max.   : 3.52950  Max.   : 2.765711
##      adherence
##  Min. :-2.94078
##  1st Qu.:-0.69257
##  Median : 0.08794
##  Mean   : 0.00000
##  3rd Qu.: 0.81430
##  Max.   : 1.77298

set.seed(1120)
km.res = kmeans(num.cov.scaled, centers = 4, nstart = 10)
km.res$centers

##      age baseline_sbp        BMI salt_intake adherence
## 1  0.26061101  -0.1715193  0.1445699 -0.3331798 -1.3480678
## 2  0.90858666   0.2343898 -0.1561342  0.6995590  0.3138607
## 3 -0.93926299   0.4109022  0.3968981  0.3220699  0.1734437
## 4 -0.07625028  -0.5837516 -0.4410942 -0.8231014  0.6876488
```

```

km.res$cluster

## [1] 4 4 3 2 3 1 1 3 3 3 3 4 4 2 4 3 1 1 2 2 2 4 1 4 2 2 3 1 3 4 3 3 2 2 4 3
## [38] 4 3 2 4 4 3 1 1 4 2 1 2 2 3 1 2 1 1 4 3 3 3 3 4 4 3 2 4 1 3 4 4 2 2 2 3 4
## [75] 4 4 4 1 4 1 1 3 3 3 4 4 4 4 4 3 3 2 2 3 4 3 1 1 1 1 2 1 2 3 4 2 1 3 3 2 1 1
## [112] 4 2 4 3 2 3 1 1 1 2 2 3 2 4 3 1 3 1 1 2 2 4 4 1 4 3 4 4 1 2 2 4 3 1 3 2 4
## [149] 2 1 4 1 1 1 2 4 4 2 2 3 1 3 3 4 2 3 2 3 1 1 3 4 1 2 3 1 3 3 3 1 2 2 4 3 3
## [186] 1 3 4 1 3 1 2 3 1 4 2 2 2 4 3 3 1 1 2 1 2 1 3 1 1 4 1 1 1 3 4 3 2 2 2 2 1
## [223] 1 2 4 4 2 3 4 1 4 2 4 2 1 4 3 3 4 1 2 4 2 1 4 2 3 3 3 3 1 1 4 1 4 2 1 1
## [260] 1 4 2 4 3 3 4 2 2 2 2 3 3 1 2 2 1 1 4 3 2 2 2 2 3 2 4 1 2 1 2 1 4 3 3 4 3
## [297] 3 3 1 3 2 4 3 2 4 3 1 2 1 3 1 1 2 3 4 4 2 2 1 1 4 2 4 2 2 4 3 2 2 4 2 3 1
## [334] 3 4 3 3 4 2 4 2 2 2 1 4 3 4 1 2 3 4 2 2 2 2 3 4 2 3 2 4 4 4 1 3 1 4 4 3
## [371] 4 4 4 1 2 3 3 4 2 2 4 1 4 4 2 3 2 4 2 1 2 3 3 3 3 3 2 3 2 1 4 4 3 1 4 1 3
## [408] 1 3 1 4 4 1 4 4 2 4 3 4 1 1 3 1 1 3 3 2 4 4 2 3 2 1 2 3 2 1 4 3 4 2 2 4 2
## [445] 4 2 2 4 2 3 3 1 4 3 1 3 4 2 2 3 4 3 2 4 2 3 1 4 2 2 2 1 3 1 3 3 2 3 2 3 2
## [482] 2 2 3 3 3 2 2 3 3 4 2 1 2 2 2 1 3 2 1 4 3 3 4 3 3 1 2 1 3 1 4 3 4 3 4 1 2
## [519] 3 4 1 3 4 3 1 2 1 3 3 4 3 2 1 1 1 1 2 3 3 1 4 1 3 2 3 1 1 2 1 3 1 2 4 3
## [556] 3 4 1 4 1 4 1 3 1 3 3 4 3 1 2 2 3 3 1 3 1 2 3 4 3 1 1 2 2 3 4 2 4 2 4 4 4
## [593] 3 3 3 2 4 3 4 4 2 2 1 4 4 1 3 4 3 1 3 2 3 3 4 3 2 3 4 1 3 3 3 4 4 1 3 4 3
## [630] 2 4 3 4 4 2 2 2 3 4 4 3 4 3 4 3 3 2 1 2 2 4 1 2 4 3 3 4 4 2 2 3 2 1 3 3 2
## [667] 2 1 3 1 3 2 1 3 4 4 2 1 1 4 2 3 2 1 2 3 4 3 1 3 1 3 4 2 1 3 4 4 3 3
```

```

J <- km.res$tot.withinss
J

## [1] 2362.18

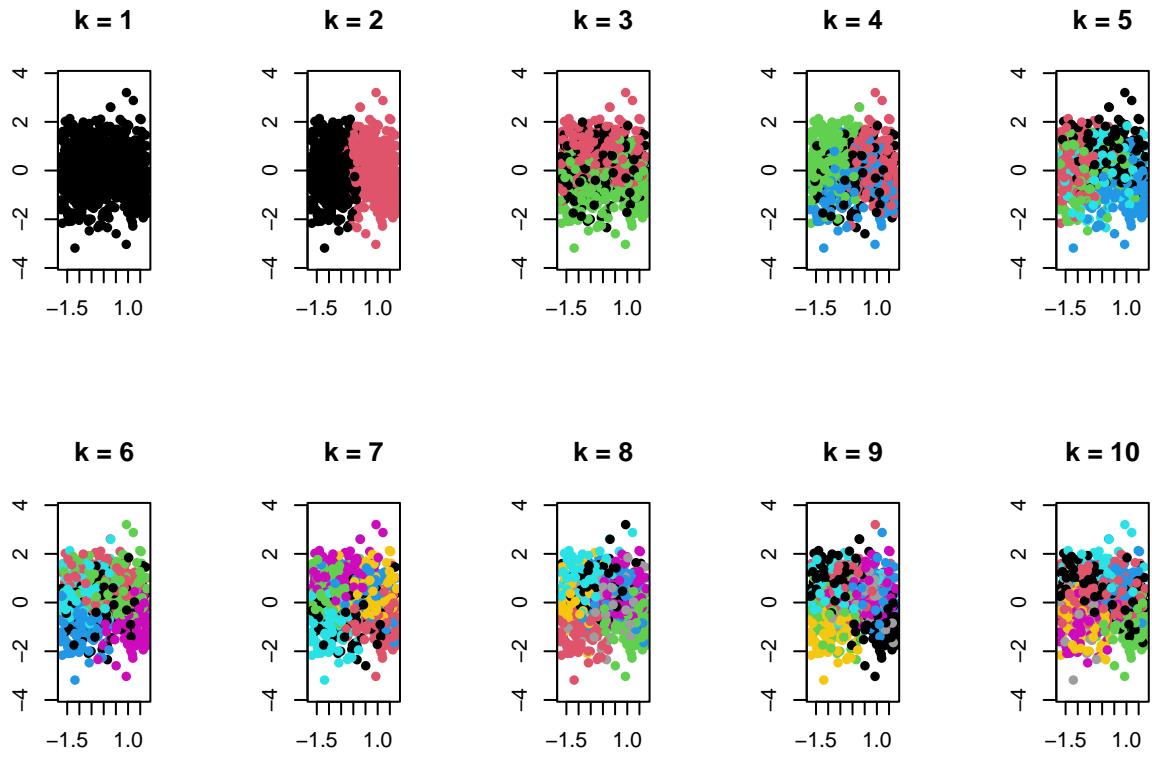
J.all = numeric(10)

layout(matrix(1:10, nrow = 2, byrow = TRUE))
for(k in 1:10){
  set.seed(1120)
  kmeans_res = kmeans(num.cov.scaled, centers = k, nstart = 10)

  plot(num.cov.scaled, pch = 16, cex = 1,
       xlab = '', ylab = '', asp = 1,
       col = kmeans_res$cluster,
       main = paste("k =", k))

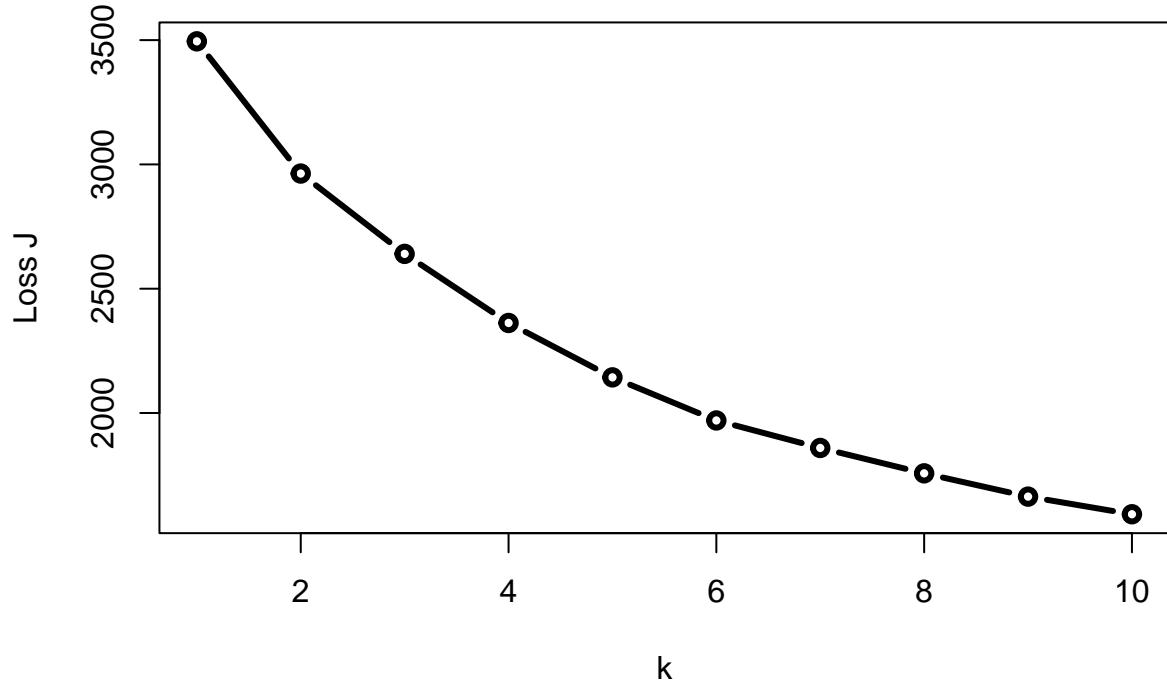
  J.all[k] = kmeans_res$tot.withinss
}

```



```
layout(1)
plot(1:10, J.all, type = 'b', lwd = 3,
      xlab = "k", ylab = "Loss J",
      main = "Loss J vs Number of Clusters")
```

## Loss J vs Number of Clusters



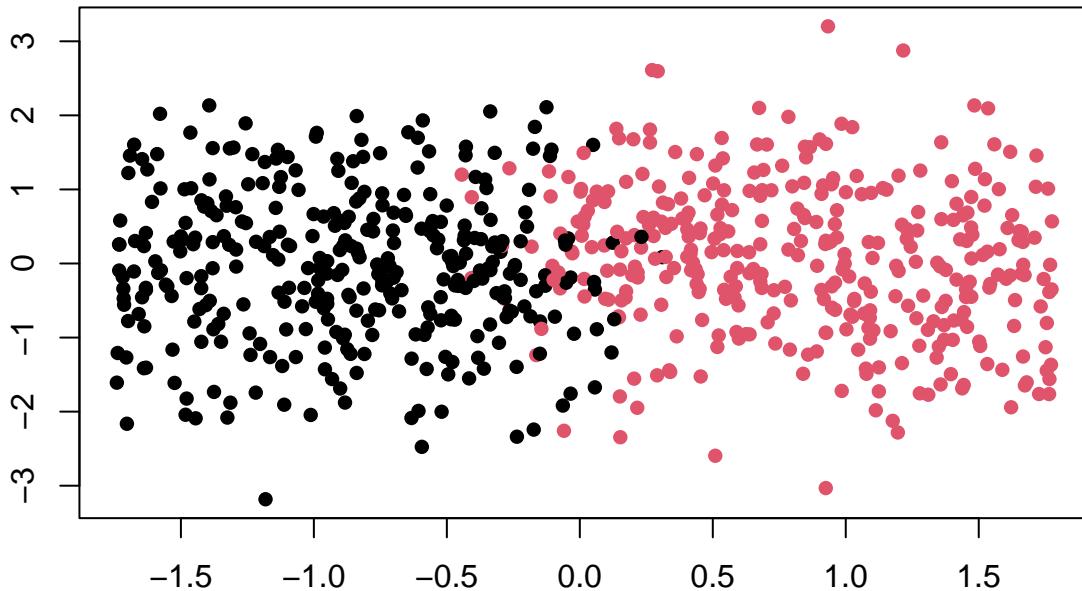
The plot does not show a clear elbow point: the decrease in the loss is smooth and gradual for all values of  $k$  from 1 to 10. When inspecting the cluster plots, none of the solutions reveals well-separated or naturally distinguishable groups, and the partitions appear mostly arbitrary. For these reasons, the data do not seem to exhibit a meaningful cluster structure, and no particular value of  $k$  is clearly supported.

Although the plot of  $J(k)$  does not show a clear elbow and the data do not seem to contain well-separated groups, if we are required to select a number of clusters with  $k > 1$ , the most reasonable choice is  $k=2$ . The decrease in the loss function is largest when moving from  $k=1$  to  $k=2$ , while further increases in  $k$  lead to smaller marginal reductions. The plot for  $k=2$  shows a mild but visible separation into two groups, whereas higher values of  $k$  produce partitions that appear mostly arbitrary. Therefore,  $k=2$  represents the most interpretable and justifiable clustering solution under the constraint  $k > 1$ .

```
set.seed(1120)
km.final <- kmeans(num.cov.scaled, centers = 2, nstart = 10)

plot(num.cov.scaled,
      pch = 16, cex = 1,
      col = km.final$cluster,
      xlab = "", ylab = "",
      main = "K-means with k = 2")
```

## K-means with k = 2



The algorithm assigns each observation to the cluster whose center (mean vector) is closest in Euclidean distance, so each group collects points that lie closer to one centroid than to the other. However, the two clusters strongly overlap, and no clear separation emerges. This is consistent with the behavior of the loss  $J(k)$ , which does not display a clear elbow. So, the solution with  $k=2$  should be seen just as a rough split of the data based on distance to the two centers, not as two real or clearly separated groups. ## Exercise 2

```
set.seed(1120)
Y_cat = factor(ifelse(data.all$Y < 17, 0, 1))
data.all.cat = data.all
data.all.cat$Y = Y_cat
n = dim(data.all)[1]
select.train = sample(1:n, n*7/10)
train.cat = data.all.cat[select.train,]
test.cat = data.all.cat[-select.train,]
summary(train.cat)
```

	## Y	age	sex	baseline_sbp	BMI
## 0:296	Min.	:30.04	Male :233	Min. : 99.41	Min. :15.99
## 1:194	1st Qu.:	42.69	Female:257	1st Qu.:137.90	1st Qu.:24.62
##	Median :	54.07		Median :148.56	Median :27.03
##	Mean :	54.44		Mean :148.66	Mean :27.13
##	3rd Qu.:	66.75		3rd Qu.:159.43	3rd Qu.:29.97
##	Max. :	79.91		Max. :189.55	Max. :40.74
##	smoker	exercise_level	salt_intake	treatment	adherence
## No :355	Low	:197	Min. : 1.124	Control:188	Min. :0.1260
## Yes:135	Moderate	:206	1st Qu.: 6.655	DrugA :154	1st Qu.:0.5376

```

##          High     : 87    Median : 7.963   DrugB   :148    Median :0.6809
##          Mean    : 7.985      Mean   :0.6631
##          3rd Qu.: 9.348      3rd Qu.:0.8166
##          Max.    :13.433      Max.   :0.9784

##  diabetes
##  No :373
##  Yes:117
##
## 
## 
## 

library(e1071)
svm.model <- svm(
  Y ~ . ,
  data   = train.cat,
  kernel = "radial",
  cost   = 10,
  gamma  = 1,
  scale   = TRUE
)

pred.train <- predict(svm.model, newdata = train.cat)
pred.test  <- predict(svm.model, newdata = test.cat)
acc.train <- mean(pred.train == train.cat$Y)
acc.test  <- mean(pred.test  == test.cat$Y)
acc.train

## [1] 1

acc.test

## [1] 0.7142857

library(caret)

## Warning: il pacchetto 'caret' è stato creato con R versione 4.4.3

## Caricamento del pacchetto richiesto: ggplot2

## Caricamento del pacchetto richiesto: lattice

set.seed(1120)
tune.radial = tune(
  svm,
  Y ~ .,
  data   = train.cat,
  kernel = "radial",
  ranges = list(
    cost  = c(0.1, 1, 10),

```

```

    gamma = c(1, 2)
  ),
  scale = TRUE
)

summary(tune.radial)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     10      1
##
## - best performance: 0.277551
##
## - Detailed performance results:
##   cost gamma      error dispersion
## 1  0.1        1  0.3959184 0.08117787
## 2  1.0        1  0.3040816 0.07723411
## 3 10.0        1  0.2775510 0.07017032
## 4  0.1        2  0.3959184 0.08117787
## 5  1.0        2  0.3979592 0.07830520
## 6 10.0        2  0.3979592 0.08120637

best.svm = tune.radial$best.model

pred.svm = predict(best.svm, newdata = test.cat)
confusionMatrix(pred.svm, test.cat$Y)

## Confusion Matrix and Statistics
##
##             Reference
## Prediction  0   1
##           0 114  54
##           1   6  36
##
##             Accuracy : 0.7143
##                 95% CI : (0.6481, 0.7743)
##       No Information Rate : 0.5714
##       P-Value [Acc > NIR] : 1.348e-05
##
##             Kappa : 0.375
##
## McNemar's Test P-Value : 1.298e-09
##
##             Sensitivity : 0.9500
##             Specificity  : 0.4000
##       Pos Pred Value : 0.6786
##       Neg Pred Value : 0.8571

```

```

##           Prevalence : 0.5714
##           Detection Rate : 0.5429
##   Detection Prevalence : 0.8000
##           Balanced Accuracy : 0.6750
##
##           'Positive' Class : 0
##

```

The best model is obtained with cost = 10 and gamma = 1, which achieves the lowest cross-validation error ( 0.278). On the test set, this SVM reaches an accuracy of about 71.4%, above the No Information Rate. Since the positive class is 0, the sensitivity is very high (0.95), meaning that most observations in class 0 are correctly classified. The specificity is much lower (0.40), showing that the model is less accurate in detecting class 1.

```

library(ROCR)

## Warning: il pacchetto 'ROCR' è stato creato con R versione 4.4.2

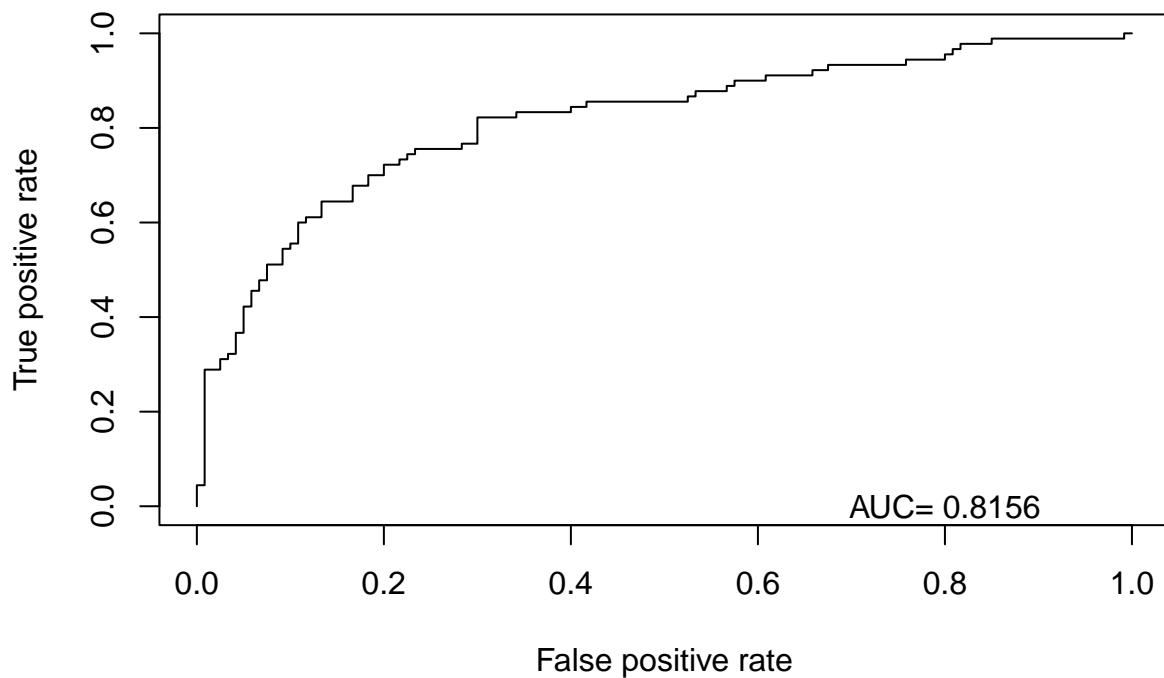
rocplot=function(pred, truth, ...){
  predob = prediction(pred, truth)
  perf = performance(predob , "tpr", "fpr")
  plot(perf ,...)
  auc = performance(predob,"auc")
  auc <- auc@y.values[[1]]
  auc = signif(auc,4)
  text(0.8,0,paste("AUC=",auc))
}

fitted.best.test = -attributes(
  predict(svm.model, test.cat, decision.values = TRUE)
)$decision.values

rocplot(fitted.best.test, test.cat$Y,
        main="ROC - Tuned SVM (test)")

```

## ROC – Tuned SVM (test)



The tuned SVM and the original SVM actually use the same hyperparameters (the tuning step confirms the initial choice), so their ROC curves and AUC are expected to be essentially the same. On the test set, the model achieves an AUC of 0.8156, indicating good discrimination between the two classes across thresholds. Overall, tuning the hyperparameters improves the model's generalization ability and leads to a clear increase in predictive accuracy. ## Exercise 3

```
n = dim(data.all)[1]
set.seed(1120)
select.train = sample(1:n,n*7/10)
train = data.all[select.train,]
test = data.all[-select.train,]

library(gam)

## Warning: il pacchetto 'gam' è stato creato con R versione 4.4.3

## Caricamento del pacchetto richiesto: splines

## Caricamento del pacchetto richiesto: foreach

## Loaded gam 1.22-6

gam.model <- gam(
  Y ~ s(age, df = 3) +
    s(baseline_sbp, df = 3) +
```

```

  s(BMI, df = 3) +
  s(salt_intake, df = 3) +
  s(adherence, df = 3) +
  sex + smoker + exercise_level + treatment + diabetes,
  data = train
)

pred.train.gam <- predict(gam.model, newdata = train)
pred.test.gam <- predict(gam.model, newdata = test)
mse.train <- mean((train$Y - pred.train.gam)^2)
mse.test <- mean((test$Y - pred.test.gam)^2)

mse.train

## [1] 121.845

mse.test

## [1] 77.18914

gam.sum <- summary(gam.model)

p <- nrow(gam.sum$parametric.anova) - 1
p.values.par <- gam.sum$parametric.anova[1:p, 5]
p.values.nonpar <- gam.sum$anova[, 3]

p.values <- c(p.values.par, p.values.nonpar)

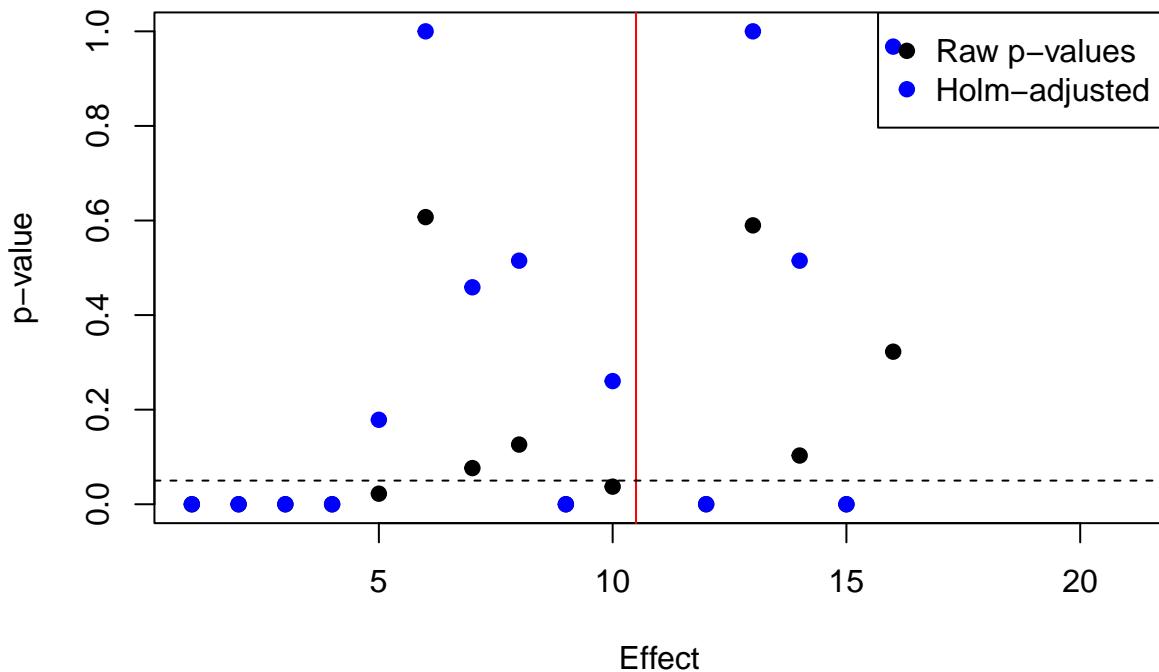
holm.p.values <- p.adjust(p.values, method = "holm")

plot(p.values, col = "black", pch = 19, ylim = c(0,1),
     xlab = "Effect", ylab = "p-value", main="Raw vs Holm-adjusted p-values")

points(holm.p.values, col = "blue", pch = 19)
abline(v = p + 0.5, col = "red")    # separates linear vs smooth
abline(h = 0.05, lty = 2)
legend("topright", legend=c("Raw p-values", "Holm-adjusted"),
       col=c("black","blue"), pch=19)

```

## Raw vs Holm-adjusted p-values



From the plot comparing raw and Holm-adjusted p-values, we observe that the Holm correction (controlling the Family-Wise Error Rate) makes the significance criterion more stringent, so fewer smooth terms remain significant at the 0.05 level. Several terms have small raw p-values, but after Holm adjustment only the strongest effects stay below the 0.05 threshold. Overall, this indicates that while the GAM initially suggests multiple potential nonlinear relationships, only the most statistically robust nonlinear effects are supported once multiple testing is taken into account.

```
summary(gam.model)

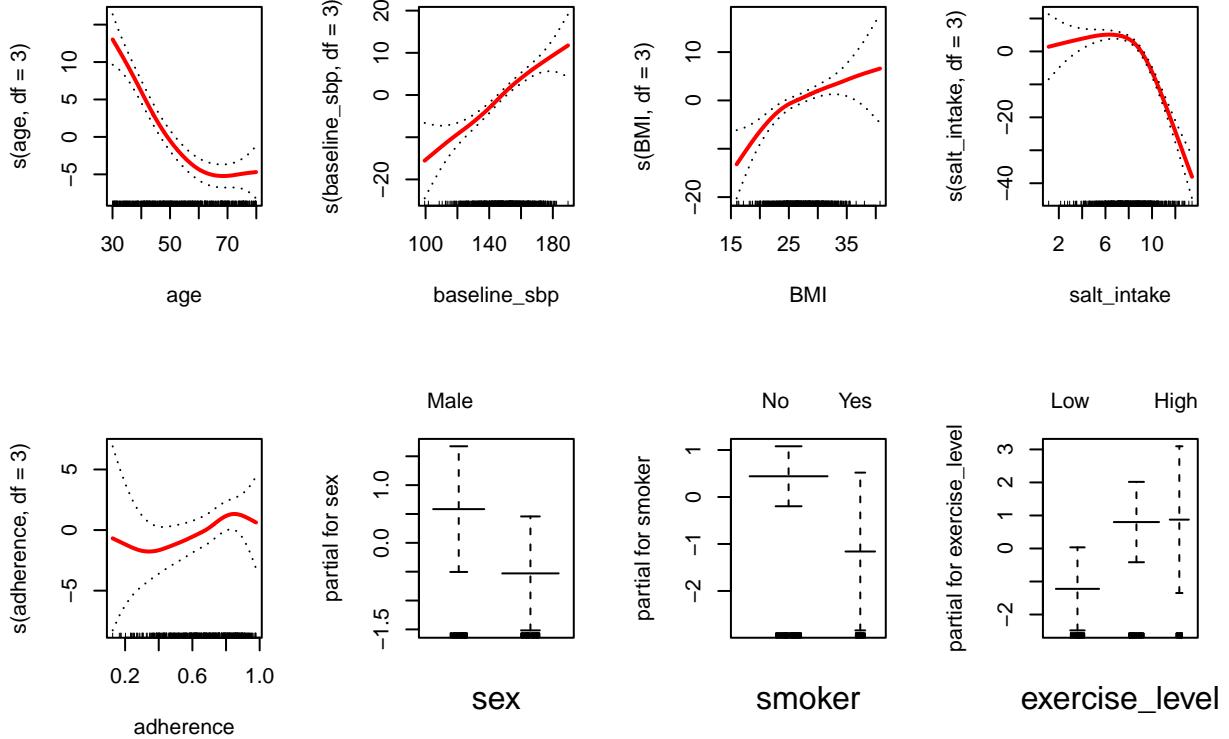
##
## Call: gam(formula = Y ~ s(age, df = 3) + s(baseline_sbp, df = 3) +
##           s(BMI, df = 3) + s(salt_intake, df = 3) + s(adherence, df = 3) +
##           sex + smoker + exercise_level + treatment + diabetes, data = train)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -108.81721  -4.71090   0.05131   5.36467  29.45625
## 
## (Dispersion Parameter for gaussian family taken to be 127.8471)
## 
## Null Deviance: 136155.9 on 489 degrees of freedom
## Residual Deviance: 59704.55 on 466.9998 degrees of freedom
## AIC: 3791.912
## 
## Number of Local Scoring Iterations: NA
## 
## Anova for Parametric Effects
```

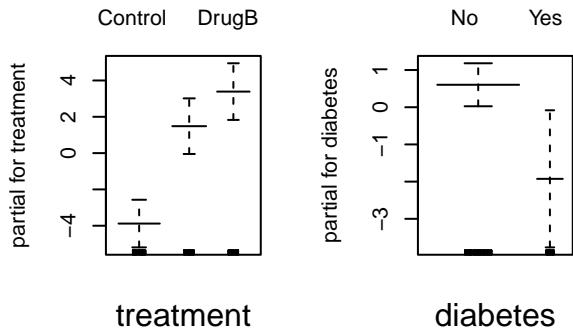
```

##                                     Df Sum Sq Mean Sq F value    Pr(>F)
## s(age, df = 3)                  1 15896 15896.3 124.3387 < 2.2e-16 ***
## s(baseline_sbp, df = 3)        1 12013 12012.7  93.9619 < 2.2e-16 ***
## s(BMI, df = 3)                 1   4750   4750.2  37.1557 2.286e-09 ***
## s(salt_intake, df = 3)         1 18855 18854.6 147.4777 < 2.2e-16 ***
## s(adherence, df = 3)           1    672    671.9   5.2556  0.02232 *
## sex                            1     34     33.8   0.2647  0.60717
## smoker                          1    403    403.0   3.1525  0.07646 .
## exercise_level                  2    532    265.8   2.0791  0.12620
## treatment                        2   4681   2340.7  18.3082 2.214e-08 ***
## diabetes                         1    558    558.2   4.3661  0.03720 *
## Residuals                       467  59705   127.8
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##                                     Npar Df Npar F      Pr(F)
## (Intercept)                      2   11.287 1.632e-05 ***
## s(age, df = 3)                   2    0.529   0.5897
## s(baseline_sbp, df = 3)          2    2.284   0.1030
## s(BMI, df = 3)                  2  57.288 < 2.2e-16 ***
## s(salt_intake, df = 3)           2   1.134   0.3226
## s(adherence, df = 3)             2
## sex
## smoker
## exercise_level
## treatment
## diabetes
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

layout(matrix(1:8,nrow=2,byrow=TRUE))
plot.Gam(gam.model , se=TRUE , col ="red",lwd=2)

```





As we can see from the smooth plots and from the ANOVA table for nonparametric effects, only age and salt\_intake show a statistically significant non-linear effect (both have  $p < 0.05$  in the nonparametric test). In contrast, the smooth terms for baseline\_sbp, BMI, and adherence are not significantly non-linear (their nonparametric p-values are  $> 0.05$ ), so their effects can be reasonably approximated as linear (or treated as weak/non-essential nonlinearities).

From the parametric part of the model and the corresponding partial-effect plots, the categorical variables sex and exercise\_level are not significant ( $p > 0.05$ ) and their partial effects largely overlap around zero. Smoker is also not significant at the 5% level (it is only borderline at the 10% level). On the other hand, treatment is clearly significant and shows a strong difference between Control and DrugB, and diabetes is significant as well, with a visible shift between levels.

```
best.gam<- gam(
  Y ~ s(age, df = 3) +
    baseline_sbp +
    BMI +
    s(salt_intake, df = 3) +
    adherence +
    treatment + diabetes,
  data = train
)

pred.train.gam <- predict(best.gam, newdata = train)
pred.test.gam <- predict(best.gam, newdata = test)
mse.train <- mean((train$Y - pred.train.gam)^2)
mse.test <- mean((test$Y - pred.test.gam)^2)
```

```

mse.train

## [1] 125.5153

mse.test

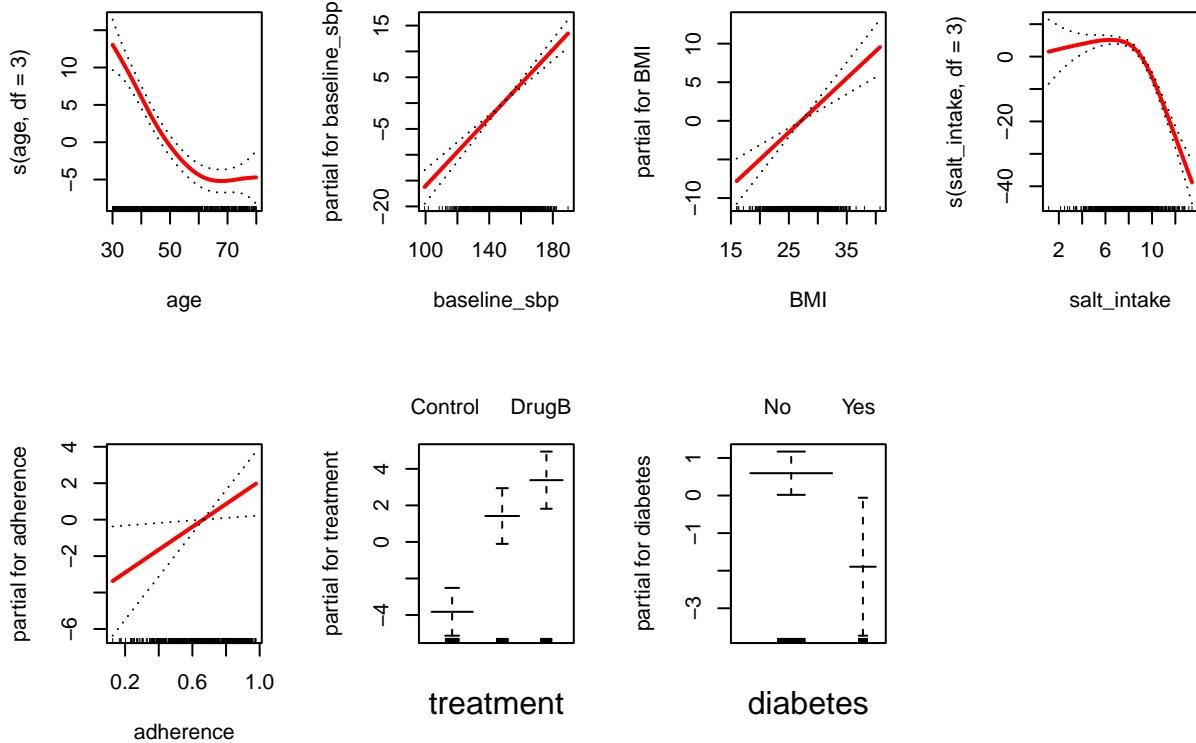
## [1] 82.36546

summary(best.gam)

##
## Call: gam(formula = Y ~ s(age, df = 3) + baseline_sbp + BMI + s(salt_intake,
##      df = 3) + adherence + treatment + diabetes, data = train)
## Deviance Residuals:
##      Min     1Q   Median     3Q    Max
## -107.2509 -4.7632  0.5207  5.6693 31.6731
##
## (Dispersion Parameter for gaussian family taken to be 128.9369)
##
## Null Deviance: 136155.9 on 489 degrees of freedom
## Residual Deviance: 61502.86 on 476.9998 degrees of freedom
## AIC: 3786.453
##
## Number of Local Scoring Iterations: NA
##
## Anova for Parametric Effects
##                               Df Sum Sq Mean Sq F value Pr(>F)
## s(age, df = 3)             1 15947 15947.1 123.6817 < 2.2e-16 ***
## baseline_sbp                1 11844 11843.9 91.8579 < 2.2e-16 ***
## BMI                          1  4714  4713.6 36.5575 2.994e-09 ***
## s(salt_intake, df = 3)      1 19003 19003.0 147.3819 < 2.2e-16 ***
## adherence                     1    764    764.5  5.9292  0.01526 *
## treatment                     2   4634   2317.0 17.9699 2.991e-08 ***
## diabetes                      1    549    549.4  4.2611  0.03954 *
## Residuals                    477 61503   128.9
##
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##                               Npar Df Npar F     Pr(F)
## (Intercept)                  2 11.455 1.383e-05 ***
## s(age, df = 3)                2 60.123 < 2.2e-16 ***
## baseline_sbp
## BMI
## s(salt_intake, df = 3)       2 60.123 < 2.2e-16 ***
## adherence
## treatment
## diabetes
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```
layout(matrix(1:8,nrow=2,byrow=TRUE))
plot.Gam(best.gam , se=TRUE , col ="red",lwd=2)
```



Age shows a clear nonlinear decreasing effect: younger patients experience a larger reduction in systolic pressure, while the improvement becomes progressively smaller with increasing age, so older patients benefit less. Salt intake has a nonlinear effect that is relatively stable at low-moderate levels and decreases only for very high daily salt consumption, suggesting that excessive salt intake reduces the effectiveness of the treatment.

In contrast, baseline\_sbp, BMI, and adherence have significant linear effects in the final model (all with  $p < 0.05$ ), while their nonlinear components are not significant, so a linear specification is appropriate for these covariates. Treatment group plays an important role: patients in the control arm achieve the smallest reduction, whereas DrugB shows a clearly larger reduction in systolic pressure. Finally, diabetes has a negative impact on the outcome: diabetic patients tend to experience a lower reduction in systolic blood pressure compared to non-diabetic individuals.

#### Exercise 4

```
set.seed(1120)
library(gbm)
```

```
## Warning: il pacchetto 'gbm' è stato creato con R versione 4.4.3
```

```
## Loaded gbm 2.2.2
```

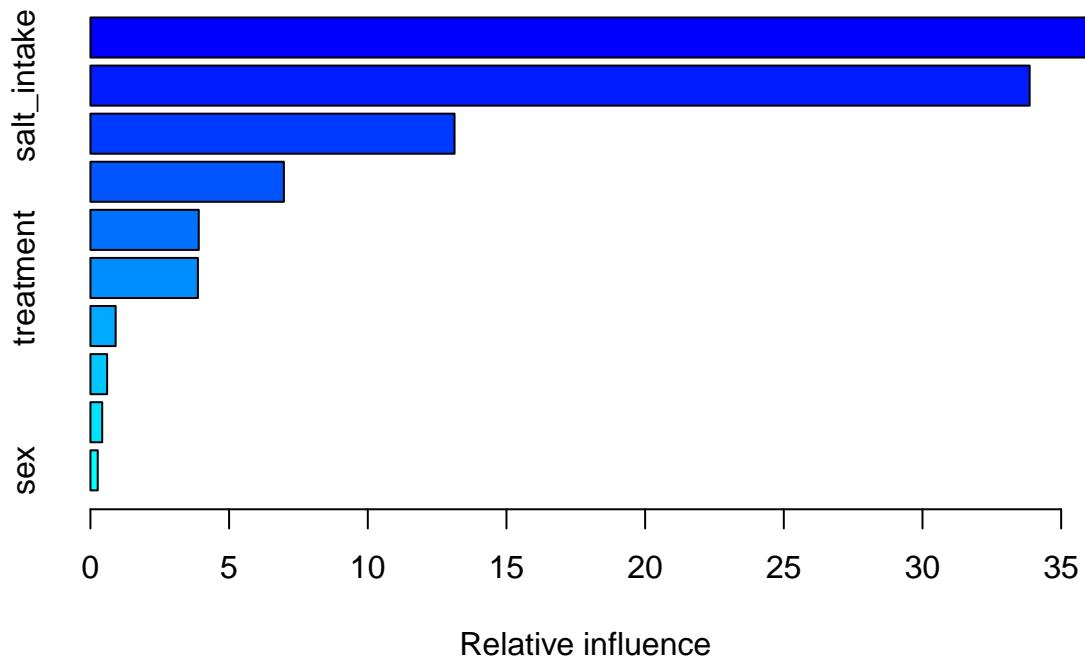
```

## This version of gbm is no longer under development. Consider transitioning to gbm3, https://github.com/tidymodels/gbm3

boost.model = gbm(
  Y ~ .,
  data          = train,
  distribution   = "gaussian",
  n.trees       = 5000,
  interaction.depth = 4,
  shrinkage     = 0.01
)

summary(boost.model)

```



```

##                                var      rel.inf
## age                         age 36.0529157
## salt_intake      salt_intake 33.8648213
## baseline_sbp    baseline_sbp 13.1278944
## BMI                          BMI  6.9743856
## adherence        adherence  3.9051428
## treatment         treatment  3.8749759
## exercise_level  exercise_level 0.9099273
## diabetes           diabetes  0.6021578
## smoker             smoker  0.4240749
## sex                  sex  0.2637043

```

```

yhat.train = predict(boost.model, newdata = train, n.trees = 5000)
MSE.train = mean((train$Y - yhat.train)^2)

yhat.test = predict(boost.model, newdata = test, n.trees = 5000)
MSE.test = mean((test$Y - yhat.test)^2)

MSE.train

## [1] 6.310776

MSE.test

## [1] 35.08872

d.explore = 1:7
params = data.frame(d = d.explore)

set.seed(1120)
indexes = sample(1:2, size = nrow(train), replace = TRUE, prob = c(0.8, 0.2))
subtrain = train[indexes == 1, ]
subtest = train[indexes == 2, ]

MSE = numeric(nrow(params))

for(ii in 1:nrow(params)){
  model.boost = gbm(
    Y ~ .,
    data = subtrain,
    distribution = "gaussian",
    n.trees = 5000,
    interaction.depth = params$d[ii],
    shrinkage = 0.01,
    verbose = FALSE
  )

  pred = predict(model.boost, newdata = subtest, n.trees = 5000)

  MSE[ii] = mean( (pred - subtest$Y)^2 )
}

opt.index = which.min(MSE)
opt.params = params[opt.index, ]
opt.params

```

```
## [1] 3
```

```
MSE_min = MSE[opt.index]
MSE_min
```

```
## [1] 37.07161
```

By evaluating different values of the interaction depth d from 1 to 7 on an internal validation split, the smallest MSE is obtained for d=3. The interaction depth controls how many interaction effects the boosting model can capture: d=1 allows only additive (non-interacting) effects, larger values of d allow interactions between predictors. An optimal value of d=3 indicates that the model benefits from including up to third-order interactions, meaning that the relationship between the predictors and the response is not purely additive but involves meaningful combinations of multiple variables.

The variable importance plot shows that age and salt intake are by far the most informative predictors, with relative influence values above 30%. These two covariates contribute the most to reducing prediction error, meaning that variation in blood-pressure reduction after 6 months is strongly driven by the patient's age and daily salt consumption. Overall, the boosting model suggests that blood-pressure reduction is mainly associated with age and salt-intake patterns, while the remaining covariates contribute substantially less to the predictive performance.

## Exercise 5

```

lm.model = lm(Y ~ ., data = train)
summary(lm.model)

##
## Call:
## lm(formula = Y ~ ., data = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -123.128  -5.418   1.242   6.723  29.717
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -10.64911  7.93549 -1.342 0.180247
## age          -0.39225  0.04135 -9.487 < 2e-16 ***
## sexFemale    -1.53321  1.17084 -1.309 0.190997
## baseline_sbp  0.28759  0.03717  7.737 6.10e-14 ***
## BMI           0.75773  0.14810  5.116 4.52e-07 ***
## smokerYes     -1.92669  1.30648 -1.475 0.140947
## exercise_levelModerate 3.20281  1.28610  2.490 0.013102 *
## exercise_levelHigh    2.65497  1.65538  1.604 0.109411
## salt_intake     -3.03806  0.29035 -10.463 < 2e-16 ***
## treatmentDrugA   5.30902  1.40917  3.767 0.000186 ***
## treatmentDrugB   7.16256  1.41035  5.079 5.46e-07 ***
## adherence        4.70773  3.14938  1.495 0.135625
## diabetesYes     -2.20343  1.36503 -1.614 0.107144
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 12.75 on 477 degrees of freedom
## Multiple R-squared:  0.4303, Adjusted R-squared:  0.4159
## F-statistic: 30.02 on 12 and 477 DF,  p-value: < 2.2e-16

yhat.test = predict(lm.model, newdata = test)
MSE.test.lm = mean((test$Y - yhat.test)^2)

MSE.test.lm

```

```

## [1] 111.7485

boost.model = gbm(
  Y ~ .,
  data      = train,
  distribution = "gaussian",
  n.trees    = 5000,
  interaction.depth = 3,
  shrinkage   = 0.01
)

yhat.train = predict(boost.model, newdata = train, n.trees = 5000)
MSE.train  = mean((train$Y - yhat.train)^2)
yhat.test  = predict(boost.model, newdata = test, n.trees = 5000)
MSE.test   = mean((test$Y - yhat.test)^2)

MSE.train

## [1] 9.589668

MSE.test

## [1] 34.76147

```

Based on the test MSE values, the boosting model is clearly the best-performing one, with an error of about 34 compared to 111 for the linear model and 82 for the GAM. This indicates that the relationship between the covariates and the reduction in systolic blood pressure is not well captured by additive (as we saw before  $d = 3$  not 1) or linear structures, while boosting successfully models the nonlinear effects and interactions in the data. The results show that age and salt intake are the strongest predictors of blood-pressure reduction, with BMI, adherence, and treatment also contributing but to a smaller extent. A further improvement could be obtained by tuning the number of trees used in the boosting model. In addition, the shrinkage parameter could also be tuned.